Harvard Business School

9-125-093
REV: APRIL 8, 2025

SURAJ SRINIVASAN
MICHELLE HU
SRIRAGHAV SRINIVASAN
RADHIKA KAK

Perplexity: Redefining Search

"Manuscripts, the printing press, libraries, search engines, and now answer engines and chatbots. This is a once in a generation moment,"

— Aravind Srinivas, CEO of Perplexity AI[1]

It was February 2025 and Aravind Srinivas, CEO and co-founder of Perplexity had just announced the release of Comet, an AI native web browser, the latest in a whirlwind of product launches from the company. He reflected on the what had driven Perplexity's meteoric growth this far. What began as a modest experiment in AI-powered search had rapidly evolved into a "knowledge discovery engine" that was increasingly defining how people searched for information in the age of generative AI. Perplexity sought to redefine search by providing quick and accurate answers to complex questions in an easy-to-digest format, powered by AI technology. The impact was undeniable: both the number of users and user queries were growing fast,[2] and some of the most prominent teachnology leaders had become vocal advocates. NVIDIA CEO Jensen Huang was one; "I have my own tutor today. And my tutor is Perplexity. I use it almost every day. It is an AI that calls upon other AIs and you could ask it all kinds of questions."[3]

The past year had brought significant wins. The firm had built its own inference infrastructure, expanded its web index, significantly improved the platform's speed and accuracy, and introduced new features, including shopping, finance, and deep research. It had expanded its loyal fan base of knowledge workers and students. The firm was valued at $9 billion in its most recent funding round in December 2024. Srinivas attributed the firm's success to its user centricity, execution velocity, and perfectly optimized AI infrastructure.

Yet, the future remained uncertain. As competition intensified, with OpenAI and Google investing in improving their own search offerings, Perplexity needed to stay ahead of technological change, and make the right strategic investments in its platform. It needed to maintain a high-quality, unbiased, and distraction-free user experience – one that had earned it the reputation of being user centric rather than advertiser focussed – while adding features to drive user engagement and adoption, and explore new monetization channels. These questions were critical for Perplexity as it vied to retain and expand

Professor Suraj Srinivasan, Michelle Hu (MBA 2025), Sriraghav Srinivasan (MBA 2025), and independent researcher Radhika Kak prepared this case. It was reviewed and approved before publication by a company designate. Funding for the development of this case was provided by Harvard Business School and not by the company. HBS cases are developed solely as the basis for class discussion. Cases are not intended to serve as endorsements, sources of primary data, or illustrations of effective or ineffective management.

its share in the highly competitive online search market, and deeply personal for Srinivas, who had traded a traditional career path to pursue his own vision of redefining search.

Company Background

Aravind Srinivas was born in Chennai, India where he completed his bachelors and master's in electrical engineering at Indian Institute of Technology (IIT), Madras. Srinivas subsequently started his PhD in computer science at UC Berkeley in 2017. There, John Schulman, co-founder of OpenAI noticed his work and invited him to intern at OpenAI. After working there and at Google's DeepMind for a few years, Srinivas felt a pull to create something of his own. In particular, his fascination with Google's search technology led him to think deeply about how AI could push the boundaries of information retrieval. Srinivas recalled the vision:

> I'm a big fan of Larry Page and Google. I've always had a drive to do something of the same scale and ambition. Being scholarly, accurate, and truthful; having the answer at your fingertips and being able to recall its sources: these are things I strive to embody.[4]

Srinivas' ambition aligned well with that of Denis Yarats, an AI researcher at NYU. Srinivas and Yarats first connected by email in 2020 after publishing nearly identical research papers on AI training methods just two days apart. This initial exchange sparked an ongoing dialogue on advancements in AI between the two. In July 2022, they teamed up, aiming to leverage large language models to create a new search experience. They were joined by two more co-founders: Johnny Ho, a former engineer at Quora, and Andrew Konwinski, a co-founder of Databricks.

The team's first product, Bird SQL, launched in December 2022. Using OpenAI's code generation tool Codex, Bird SQL transformed natural language prompts into code to search databases like Twitter. This early prototype—a searchable chat interface—quickly gained traction on Twitter, even catching the attention of Twitter co-founder Jack Dorsey. Bird SQL could surface insights such as "people Sam Altman engages with the most" or "top posts about ChatGPT," sparking popular interest.

Despite this success, early chatbots and language models were limited by their reliance on historical data, lacking real-time capabilities. Bird SQL had a limited shelf life as a single-purpose chatbot. The team faced a crucial early decision: whether to pursue vertical-specific applications or build a general-purpose search engine. The founders settled on the latter, envisioning a platform that prioritized up-to-date information with reliable sourcing and citations. They wanted to build a wide-reaching search engine that combined a web index with a chatbot style interface. Srinivas reflected:

> For two decades, we've been conditioned to use keywords to search the web because that's how the leading search engines were designed. The traditional approach to search was to return ten blue links, which users then had to comb through to find the information they were looking for. In recent years, this model has been evolving to provide direct answers to users' questions. LLM's are changing the way that we interact with computers to both find and consume information. That's what I mean by "answer engine." Our goal is to answer more complex questions that require synthesizing content from multiple pages and providing fast accurate answers using LLMs.[5][6]

This decision laid the foundation for what would eventually become Perplexity.

The First Prototype and Early Days

Perplexity began its evolution into a general search product in late 2022. The team built upon existing infrastructure, initially combining Bing API with GPT-3.5 to launch the first version of the product. Srinivas explained, "The decision to not invest in pre-training and massive GPU infrastructure at the start wasn't easy. We focused on delivering a great product, even if it meant being called a wrapper. Internally, we had strong belief that wrapping all human knowledge is still incredibly valuable."

The early version of Perplexity started as a Slack bot in October 2022, targeting specific use cases like answering employee medical insurance queries in "answer engine" format unlike other conversational chatbots designed for open-ended interactions. This initial prototype provided valuable insights into user needs and technical feasibility, and highlighted significant technical challenges, with initial query latency at seven seconds and frequent crashes due to rate limit issues.

To scale the system from niche use cases to general search, the team needed to overhaul its approach. At the core, it needed to tackle three key issues: speed optimization, index quality improvement, and custom infrastructure development. On speed, the goal was to become one of the fastest chatbot services. Yarats explained, "Latency matters. If you're slow, users leave. But speed without quality means nothing." On index quality, unlike competitors who were focused on broadly expanding their index, Perplexity decided to prioritize domains relevant to knowledge discovery. On infrastructure development, Perplexity invested in creating its inference infrastructure using open-source models, ensuring scalability and independence from third-party limitations.

AI Orchestration: The Core of Perplexity's Moat

Perpexity's team gradually built out an AI system with orchestration designed to address the twin issues of providing reliable, accurate answers while maintaining speed and efficiency. Yarats explained, "It's not just about stitching systems together—it's about orchestrating every part of the stack to create answers, not links."

This system constituted four critical components:

- **The Reformulator:** This module played a pivotal role in query understanding by breaking down complex or ambiguous user queries into smaller, actionable sub-queries. By leveraging large language models' reasoning capabilities, the Reformulator ensured that nuanced queries were effectively interpreted, even when user intent was unclear or multifaceted. Yarats emphasized, "Query understanding isn't just about keywords—it's about breaking down ambiguity into something actionable."

- **The Index:** Perplexity built its own index, mapping the internet into a vector database with a focus on high-quality sources. It used LLM-based analysis to assess the trustworthiness and relevance of content, bypassing traditional click signals. Yarats explained; "We don't need clicks to infer relevance—LLMs do that better by analyzing the text directly."

- **Retrieval-Augmented Generation (RAG):** This technique ensured the accuracy of responses by grounding them in the retrieved information. The system extracted and ranked content snippets from indexed sources before synthesizing them into coherent answers.

- **Large Language Models (LLMs):** Perplexity combined in-house and external open-source LLMs, including OpenAI's GPT-4, Meta's Llama 2, and Anthropic's Claude, to seamlessly

integrate retrieval, synthesis, and reasoning, ensuring that responses were not only fast, but also contextually precise and clear.

Building a Unique Index

Rather than focus on index size, Perplexity prioritized index quality, specifically targeting domains relevant to AI chatbots and knowledge workers. The company developed a content ranking system that found relevant documents and extracted useful snippets of content. This process involved LLM-based analysis of content quality and reasoning, enabling the platform to determine document relevance without relying on traditional click signals, assess content trustworthiness through AI-driven analysis, and extract pertinent paragraphs for answer synthesis. For example, sites like the New York Times were generally given higher importance than Substack posts. The tech capabilities were rooted in sophisticated crawling and indexing capabilities with custom parsers. Srinivas noted, "The accuracy of our platform has been consistently improving, primarily because we keep expanding and improving the quality of the index."

User-Centric Query Design and Query Reformulation

Perplexity aimed to make knowledge discovery as accessible and user-friendly as possible. The platform's interface was designed to make the search experience feel like interacting with a knowledge assistant, making it easy for users to explore nuanced questions and topics in-depth. Srinivas took inspiration from Google and Apple with regards to UI design, opting to minimize the number of options for users to choose from in any given interaction, with the belief that interface design should be intuitive and self-explanatory. He explained:

> Our bias is toward removing things than adding them. People sometimes complain that Google doesn't add enough new features, yet they may not realize that doing so creates more responsibility and confusion for users. We like to remove as much friction as possible.[7]

The team introduced features like "Your Threads" to save search interactions, and a Copilot feature leveraging GPT-4 for interactive assistance, to help users narrow down on a query by asking clarifying questions. Upon introduction of a new feature, the team closely followed interaction patterns. If a feature was not used, it was removed.[8]

Perplexity had strategically chosen to avoid monetization through ads. Srinivas reflected:

> If I were an advertiser, I would focus on describing my product as accurately as possible so that an LLM considers it citation-worthy. Instead of optimizing for clicks, I would optimize for high-quality content. Some journals will publish your article if you pay, but their reputation is much lower than peer-reviewed journals. I'm still trying to wrap my head around what a sponsored citation means. In the next iteration of the internet, clear, trustworthy sources of information will ideally only cite content that other, verified sources have provided. In this scenario, the LLM will be the judge.[9]

The founding team believed that keeping the platform ad-free enabled Perplexity to offer users a better and more distraction-free experience compared with Google and other ad-driven search sites. Srinivas reflected: "We realized that competing with Google would be manageable because they have a vested interest in protecting their ad revenue, which limits their ability to provide direct answers. If Google focused solely on providing correct answers, it could dramatically improve its search product with minimal effort. It's business interests that prevent them from doing so."[10] He later added, "What

makes me confident is the fact that, if they want to do better than us, they would basically have to kill their own business model."[11]

To improve ease of use, Perplexity was working on reducing the need for prompt engineering, cognizant that creating perfect prompts was a significant barrier to use for users less familiar with AI technology. The team developed a sophisticated "Reformulator" which transformed and broke down user queries into multiple related concepts and questions to enable the model to best understand and respond to the user's original intent. The focus on improving the platform's understanding of user intent went beyond just query reformulation. Perplexity employed advanced natural language processing techniques to analyze the context and nuances of user queries. This contextual understanding allowed the system to provide more relevant and accurate responses, even when the initial query was ambiguous or incomplete. Additionally, heuristics and data-driven learning from past-queries were used to improve search results. Srinivas explained, "Don't blame users for not phrasing their questions correctly. Its our job to fix these things and make the product more intuitive."[12]

Srinivas strongly believed that Perplexity's easy-to-use interface was key to its success, explaining:

> Bing tried to incorporate multiple features into one product, such as search, chat, and multi-turn conversations using GPT-4, which cluttered the user experience. By contrast, we focused only on creating an answer engine with citations and avoided freeform conversations. This clear focus helped us create a useful product, while Bing's product, despite the hype was less clear about its exact use cases.[13]

Balancing Accuracy and Latency

A critical challenge for Perplexity was optimizing for speed while maintaining a certain minimum standard for the quality of responses. Yarats explained, "There's no point in being fast if you're wrong, but accuracy without speed is equally useless." Maintaining the right balance between these two critical variables was central to user satisfaction and retention. The team continuously refined its models to enhance their precision and reliability, ensuring responses were not only factually correct but also contextually relevant. Simultaneously, they focused on optimizing the underlying infrastructure and algorithms to minimize latency. This included making advancements in query processing, dynamic caching, and inference efficiency, which allowed the system to deliver high-quality results with minimal delays.

Building Infrastructure to Handle Scale

Perplexity invested in developing robust in-house infrastructure to support its expansion to ensure scalability and maintain control over critical operations while minimizing dependency on external providers. Indeed, as user demands increased, Perplexity's infrastructure had to accommodate higher traffic volumes, more complex queries, and the integration of new features without compromising on speed or reliability. By building proprietary training pipelines and inference systems, Perplexity established a foundation that enabled the efficient training and deployment of its models while providing the flexibility to optimize for performance and cost-effectiveness at scale. The team continuously created new avenues to handle spikes in demand, deploy updates seamlessly, and maintain consistent performance across geographies and use cases.

Driving Growth

Perplexity's initial focus was on building early user engagement and retention on the platform, tracked via growth of users and queries per user. Srinivas explained, "A big thing on my mind is making people aware that we exist and growing our user base. Earning users' trust so they feel comfortable making the switch [from an incumbent platform] is a significant challenge."[14] The firm hired Dmitry Shevelenko as Chief Business Officer in October 2023 to spearhead partnerships, growth strategy, and go-to-market efforts, with the goal of expanding Perplexity's footprint. He brought extensive experience from his leadership roles at Uber, Facebook (now Meta), and LinkedIn, adding expertise in strategic partnerships, product management, and scaling technology platforms to Perplexity's executive team.

Under Shevelenko, Perplexity focused greater attention on garnering consumer impressions and brand visibility. Through an active presence in social media, in industry conversations, and regular user engagement, Perplexity was building up a core user group of advanced users. Over 80% of the platform's users had an undergraduate degree, 65% were employed in high-income white-collar jobs including medicine, law, and software, and 30% were in senior leadership positions.[15] In the summer of 2023, Perplexity introduced a $20/month premium subscription model called Perplexity Pro for high-demand users that demanded a high-quality search experience, offering users access to advanced problem-solving, multi-step reasoning, and premium data sources.

Rapid product development marked a key component of Perplexity's growth strategy, with the goal staying top-of-mind for consumers who used it and generated continuous buzz. Srinivas explained that the firm's agility in introducing new user-centric features was central to its success, explaining, "We're comfortable rolling out products that aren't fully perfect yet and iterating based on real user feedback. That's where we have the edge over big players like Google. Momentum is everything. Shipping fast, learning fast, and staying laser-focused on user growth are at the core of what we do." Dmitry added, "Part of what keeps people engaged and excited about Perplexity is that they feel like it's always getting better. There's always something new or enhanced about the experience."

Partnerships, too, played a crucial role in Perplexity's growth.The team forged strategic alliances with telecom companies globally, offering Perplexity Pro to their subscribers and integrating its API into apps and services. These partnerships spanned countries like Germany, South Korea, Japan, the UK, Spain, and most recently, India, solidifying Perplexity's presence in key global markets and bringing in a host of new users. Perplexity also launched partnerships with enterprises like Uber, Samsung Next, Yelp, and LinkedIn to establish new distribution channels and avenues to reach users. Dmitry explained: "It becomes a much more authentic form of marketing when a trusted brand that you already rely on and respect is telling you to use Perplexity."

2024 was marked by several pivotal product launches and strategic partnerships. In April, the company unveiled *Enterprise Pro*, its first B2B solution, featuring advanced security and control capabilities, including solutions that enabled integrating private, public, and third-party datasets, specifically tailored to meet enterprise needs.[16] This was driven by the firm's desire to diversify its growth and revenue channels. Shevelenko explained, "We want knowledge workers to be able to use Perplexity at work and at home." The team hoped that its combination of a consumer-friendly interface with tailored solutions geared to meet enterprise-specific needs would enable it to differentiate itself from other enterprise search solutions.

In May, it unveiled *Perplexity Pages*, a tool that enabled users to create and publish detailed, topic-specific articles to showcase their knowledge, interest, or expertise on a range of subjects, from personal

hobbies and professional insights to academic research.[17] The goal was to foster a sense of ownership and community amongst the platform's user base of learners and knowledge seekers.

In July, it launched *Pro Search*, a tool that introduced powerful problem-solving functionalities, including multi-step reasoning and enhanced capabilities for math and programming. Later that month, in response to legal challenges from publishers on content usage and intellectual property rights, Perplexity rolled out the *Publishers' Program*, offering leading media organizations like *TIME*, *Der Spiegel*, and *The Texas Tribune* benefits including revenue-sharing opportunities, API access, and complimentary *Enterprise Pro* subscriptions for their employees, in acknowledgment of their importance to the platform.[18] Srinivas conceded that Perplexity's platform likely sent less traffic to publisher websites than traditional search, but also added that the traffic which was directed was of a higher quality.[19] Shevelenko emphasized: "Publishers' businesses have been disrupted every time there's new tech. We tried to take a long-term view by committing to revenue sharing, ensuring that we support the ecosystem rather than undermine it."

To drive greater platform adoption amongst the youth, Perplexity extended *Enterprise Pro* to schools and nonprofit organizations at a reduced cost, enabling its integration into students' learning experiences and educators' workflows.[20] Perplexity also launched a promotion for university students, offering them a free month of *Perplexity Pro*. Campuses that achieved 500 sign-ups before September 15th unlocked a special upgrade, with all participating students granted a full year of free access to *Perplexity Pro*.

In February 2025, Perplexity released a *Deep Research* feature, a function to autonomously execute search, read through hundreds of sources, and reason through the material to deliver a synthesized comprehensive report on a topic. It used a framework called test time compute which mimicked human cognitive capabilities, and could identify patterns, reconcile contradictory information, prioritize credible sources, and iteratively improve its own understanding through repetitive analysis. The tool was offered as a steep discount compared with similar tools launched by Google and OpenAI, and was promoted as being especially useful for use in finance, marketing, technology, healthcare, and travel. Building on this momentum, Perplexity extended Deep Research into enterprise and developer ecosystems. Deep Research for the Enterprise enabled businesses to query both internal and external sources with enterprise-grade security. Meanwhile, the launch of the Sonar API allowed developers to build custom AI-powered research agents and workflows, reinforcing Perplexity's expansion beyond consumer search into AI-driven enterprise and developer solutions.

The company also most recently announced Comet browser, aimed at integrating AI-native search directly into a seamless browsing experience, to further differentiate its platform from traditional search engines. Additionally, Perplexity Voice Mode was introduced as a next-generation search interface, enabling users to engage in real-time, conversational interactions powered by multilingual, context-aware AI. By incorporating real-time voice and real-time search, Perplexity positioned itself at the forefront of AI-driven multimodal search experiences, challenging the long-standing paradigm of text-based search engines. Meanwhile, the launch of Perplexity's first open-weights model, R1 1776, was aimed at showing its commitment to transparency and unbiased AI.

This progress was showing results, and key performance indicators (KPIs) including the number of users, number of queries per user, user engagement, and retention rates, were all on an upward trajectory, signaling robust market penetration in the consumer and enterprise segments. The platform had 15 million active users, up from the 10 million at the start of the year, and around 50 million global visitors each month. Around 2 million people visited the site daily.[21] Query volumes had grown from 530 million in 2023 to more than 1 billion queries in the first half of 2024, and more than 400 million

queries a month as per most recent figures.[22] Most users were in the 18 – 34 age range and from the US (22.5% of total traffic), followed by India (8.36%) and then Indonesia (6.10%). Some estimates suggested that the platform was set to reach annual recurring revenues of $35 million, up sharply from $10 million in 2023, and $5 million in 2022.[23]

The Future: Challenges and Decisions

By early 2025, Perplexity had carved out a niche for itself as go-to search platform within a certain segment of knowledge-seeking users. It's value proposition built on its AI orchestration which combined a custom-built index and continuous optimization to balance accuracy and latency, and a user-centric approach, enabled it to continue making inroads with users. Still, redefining search was an ambitious mission in the online search industry where Google still held a dominant 90% market share.[24] Srinivas reflected:

> For the moment, it seems like the world is still happy with Google, as their traffic has not materially changed. Just as Google and Facebook transformed how people consume news, a shift away from traditional search engines will happen eventually. Yet, until our product achieves full parity with Google search, and then gets 10x better, people will stick with the status quo. It might take a few months or a year to reach that point. Steady, mindful progress will help us become part of people's day-to-day workflows.[25]

He added, "Some people think we're stupid to even try to compete with Google. After all, search is their crown jewel. But here's the core insight we stumbled upon: if people stop clicking on links, the ad economy dies. Google's search exists in service of its ads, not the other way around."

The firm's leadership contemplated its future direction amid a fast evolving industry landscape.

Individual users constituted the central and most integral component of Perplexity's platform; if the firm was to realize its founding mission of redefining search, it needed to grow its user base while increasing customer stickiness. Perplexity's ultimate vision was to develop an end-to-end solution for user needs, from search, answers to queries, the possibility to purchase things, and a task assistant. This would require sophistication in AI orchestration at the backend, but ease of use at the customer end. As the platform added new features, including sports, travel, entertainment, and flights, with this goal in mind, it needed to seamlessly integrate new offerings into a single unified experience so as to maintain a high-quality user experience, and ensure that the firm's early adopters who liked its knowledge-focused UI, stayed. Srinivas explained: "Magical consumer products just work. We need to be obsessed with the consumer to get things right."

One question with respect to maintaining a high-quality and distraction-free user search experience was on whether and how to integrate advertisements as a component of the platform's monetization strategy. One option could be to integrate advertising with a shopping experience. However, the platform's promise of an unbiased search experience had so far been a core source of differentiation, one that the founders themselves had vocally promoted. Striking this balance required alignment between the interests of brands and the expectations of consumers. Perplexity would need to introduce ads in a way that preserved user trust and reinforced the platform's core value proposition. Shevelenko reflected:

> We have two principles for any experimentation we're doing with advertising. First, never change the answer based on who the advertiser is. The integrity of the user experience is paramount. Second, don't make the user experience suffer. A big part of why people

prefer Perplexity over other search engines is the clutter-free interface without intrusive sponsored answers.

Simultaneously, the team wondered how big it should bet on the enterprise segment. Perplexity believed that its consumer-friendly interface, combined with robust enterprise features like data privacy and integration capabilities, positioned the platform uniquely to grow its enterprise user base. Growing this segment was financially lucrative for the firm given enterprise users' higher willingness to pay, and higher enterprise adoption also fueled greater consumer uptake, and vice-versa, creating a virtuous cycle. Still, competing in the enterprise search space alongside established platforms and emerging AI-native competitors that offered sophisticated semantic search and knowledge retrieval capabilities was challenging. Given enterprise users' concerns over data privacy, prompt handling, and hallucinations, competing in this segment would require significant investment in infrastructure, product development, and customer support and education initiatives.

Additionally, the team stayed vigilant about future threats to the business, of which rapid technological change was arguably the biggest. Staying agile was imperative. Perplexity's strategy was focused on making quick decisions and experimenting rapidly without over-optimizing too early. Still, as the cost of LLM inference continued to decline, it would be easier for new-entrants to launch application-layer search models, while incumbents like Google continued to invest. Perplexity needed to stay at the very forefront of product development, innovation, and user centricity to avoid being out-competed. For example, frontier developments in AI were enabling new applications of agentic AI, opening up opportunities in assistant services, including the facilitation of knowledge discovery and action-taking on behalf of the user. Perplexity needed to stay ahead of such trends. Srinivas stressed the urgency of the situation and the need to continuously take action: "The biggest risk is slowing down. We need to execute and ship faster, not get distracted by competitors. If we're not pushing the system fast enough, we're not swinging big enough."

As Perplexity expanded its feature set, a key strategic question emerged: how should it prioritize development without succumbing to feature sprawl? Perplexity now offered deep research, reasoning, internal and external search, access to multiple AI models, and collaborative knowledge-sharing through Spaces—all within a single platform. While this breadth of capabilities reinforced its differentiation, it also posed a risk: could Perplexity maintain a coherent, high-quality user experience while building in many directions?

With each new feature, the company needed to ensure that the sum was greater than the parts—that these capabilities worked together seamlessly to enhance user engagement rather than overwhelm or dilute the experience. Should the company double down on deep research and enterprise integrations, or prioritize voice search as the next frontier? Could it continue to win users across knowledge workers, casual consumers, and enterprises alike, or would its growing scope necessitate sharper segmentation? The answers to these questions would be critical to ensuring that Perplexity could maintain its growth trajectory.

Exhibit 1 Leadership Bios

Aravind Srinivas is the CEO of Perplexity, the conversational "answer engine" that provides precise, user-focused answers to queries — with in-line citations. Aravind co-founded the company in 2022 after working as a research scientist at OpenAI, Google, and DeepMind. To date, Perplexity has raised over $165 million from investors including Jeff Bezos, Nat Friedman, Elad Gil, NVIDIA, and the late Susan Wojciki. He has a PhD in computer science from UC Berkeley and a Bachelors and Masters in Electrical Engineering from the Indian Institute of Technology, Madras.

Denis Yarats is the Co-founder and CTO of Perplexity, the conversational "answer engine" that provides precise, user-focused answers to queries — with in-line citations. Denis co-founded the company in 2022 after working as a research scientist at Facebook. He has a PhD in Computer Science from NYU.

Johnny Ho is the co-founder and Chief Strategy Officer of Perplexity, where he leads product. Before founding Perplexity, he worked as a high-frequency trader at Tower Research Capital and as an engineer at Quora. A former competitive programmer, Johnny is a three-time gold medalist at the International Olympiad in Informatics (IOI). In 2012, he ranked first in the world at IOI, earning a perfect score. He holds a degree in math and computer science from Harvard University and is based in New York City.

Dmitry Shevelenko is the Chief Business Officer of Perplexity, where he leads all aspects of the company's business growth, including strategic consumer and enterprise partnerships. Previously, he was the founder of Tortoise, a retail and commerce automation startup. He has also worked in business development and product leadership at Uber, LinkedIn, and Meta. He has a BA in Anthropology from Columbia University.

Source: Casewriters

Exhibit 2 Perplexity's Angel Investors

Source: Caompany Documents

Exhibit 3 Perplexity's Quarterly Query Growth (2023-2024)

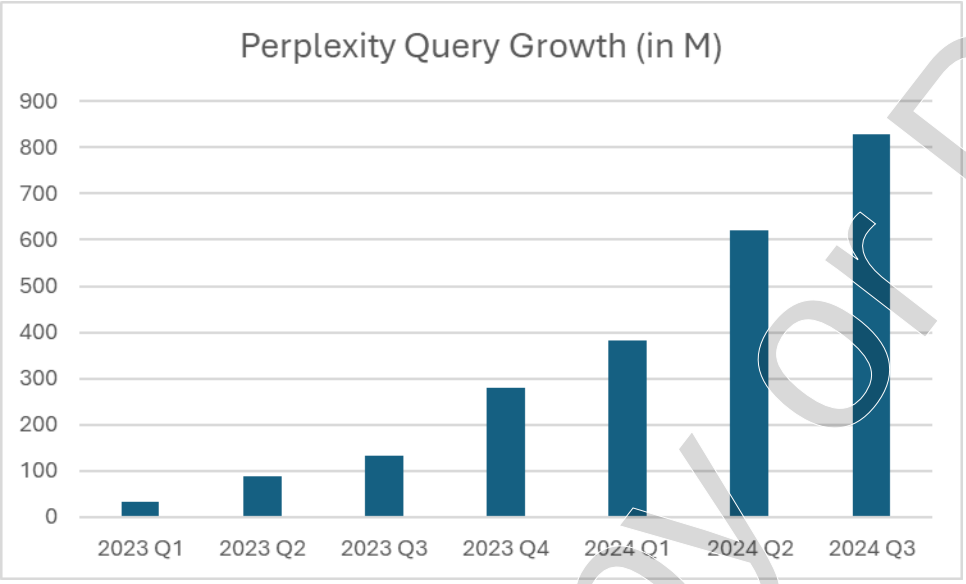

Source: Company Documents

Exhibit 4 Perplexity Spaces Dashboard

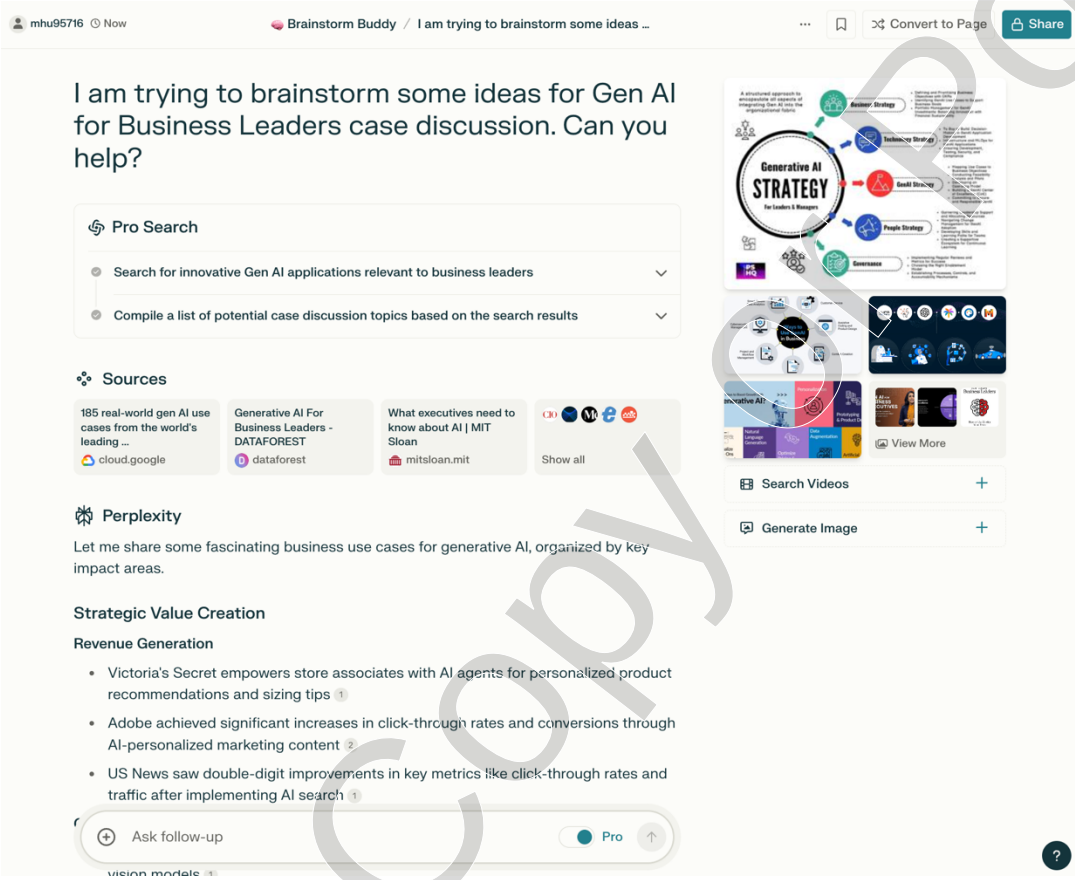

Source: Company Documents

Exhibit 5 Perplexity Finance Dashboard

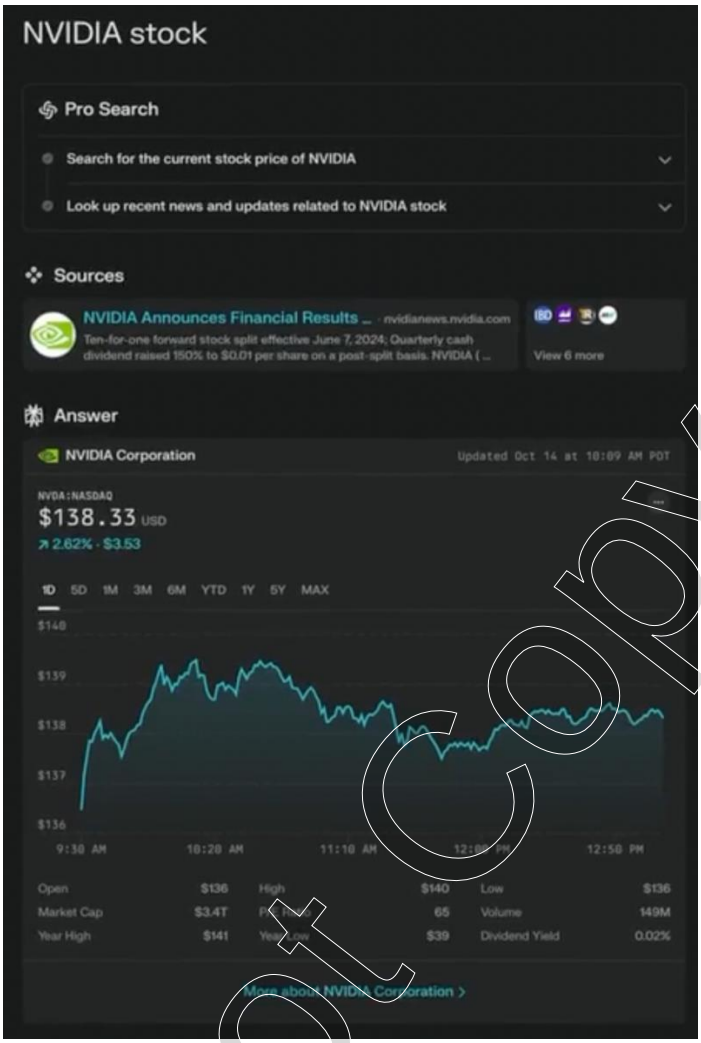

Source: Company Documents

Endnotes

[1] Joanne Chen, "How Perplexity.ai Is Pioneering the Future of Search," Forbes, September 6, 2023, https://www.forbes.com/sites/joannechen/2023/09/06/how-perplexityai-is-pioneering-the-future-of-search/, accessed March 2025

[2] Perplexity internal documents

[3] "Nvidia CEO on Artificial Intelligence and Energy." C-Span, 27 Sept. 2024, www.c-span.org/video/?538744-1%2Fnvidia-ceo-artificial-intelligence-energy, accessed March 2025

[4] Joanne Chen, "How Perplexity.ai Is Pioneering the Future of Search," Forbes, September 6, 2023, https://www.forbes.com/sites/joannechen/2023/09/06/how-perplexityai-is-pioneering-the-future-of-search/, accessed March 2025

[5] Ibid

[6] Ibid

[7] Ibid

[8] Ibid

[9] Ibid

[10] Ibid

[11] Kevin Roose, "Can This AI-Powered Search Engine Replace Google? It Has for Me," New York Times, February 1, 2024, https://www.nytimes.com/2024/02/01/technology/perplexity-search-ai-google.html, accessed March 2025

[12] Joanne Chen, "How Perplexity.ai Is Pioneering the Future of Search," Forbes, September 6, 2023, https://www.forbes.com/sites/joannechen/2023/09/06/how-perplexityai-is-pioneering-the-future-of-search/, accessed March 2025

[13] Ibid

[14] Joanne Chen, "How Perplexity.ai Is Pioneering the Future of Search," Forbes, September 6, 2023, https://www.forbes.com/sites/joannechen/2023/09/06/how-perplexityai-is-pioneering-the-future-of-search/, accessed March 2025

[15] Alexander Procter, "How Perplexity AI's Bold Advertising Move is Shaking up AI Searching," Okoone Spark, August 29, 2024, https://www.okoone.com/spark/marketing-growth/how-perplexity-ais-bold-advertising-move-is-shaking-up-ai-searching/, accessed March 2025

[16] "Perplexity Launches Enterprise Pro." Perplexity, Apr 23, 2024, www.perplexity.ai/hub/blog/perplexity-launches-enterprise-pro, accessed March 2025

[17] "Introducing Perplexity Pages." Perplexity, 30 May 2024, www.perplexity.ai/hub/blog/perplexity-pages, accessed March 2025

[18] "Introducing the Perplexity Publishers' Program." Perplexity, July 30, 2024, www.perplexity.ai/hub/blog/introducing-the-perplexity-publishers-program., accessed March 2025

[19] Kevin Roose, "Can This AI-Powered Search Engine Replace Google? It Has for Me," New York Times, February 1, 2024, https://www.nytimes.com/2024/02/01/technology/perplexity-search-ai-google.html, accessed March 2025

[20] "Bringing Perplexity to Education and Not-for-Profits." Perplexity, June 27, 2024, www.perplexity.ai/hub/blog/bringing-perplexity-to-education-and-not-for-profits, accessed March 2025

[21] Naveen Kumar, "Perplexity AI Statiwstics (2025) – Users and Revenue," Demand Sage, October 10, 2024, https://www.demandsage.com/perplexity-ai-statistics/, accessed March 2025

[22] Perplexity internal documents, Exhibit 3, User Query Growth

[23] Naveen Kumar, "Perplexity AI Statiwstics (2025) – Users and Revenue," Demand Sage, October 10, 2024, https://www.demandsage.com/perplexity-ai-statistics/, accessed March 2025

[24] GS Stat Counteer, "Search Engine Market Share Worldwide: March 2024 – March 2025," https://gs.statcounter.com/search-engine-market-share, accessed March 2025

[25] Joanne Chen, "How Perplexity.ai Is Pioneering the Future of Search," Forbes, September 6, 2023, https://www.forbes.com/sites/joannechen/2023/09/06/how-perplexityai-is-pioneering-the-future-of-search/, accessed March 2025